

Zinan Lin

Microsoft Building 99, 14820 NE 36th St – Redmond, WA 98052 – USA

✉ linzinan1995@gmail.com • 🌐 zinanlin.me

🔍 scholar.google.com/citations?user=67nE-wQ_g_cC

🔗 github.com/fjxmlzn • 🐦 [lin_zinan](https://twitter.com/lin_zinan)

Employment

Microsoft Research

Principal Researcher

Redmond, WA, USA

Mar. 2026-Present

Microsoft Research

Senior Researcher

Redmond, WA, USA

Oct. 2022-Mar. 2026

Education

Carnegie Mellon University

Ph.D., Department of Electrical and Computer Engineering

Secondary Master in Machine Learning, School of Computer Science

Secondary Master in Electrical and Computer Engineering

Ph.D. advisors: Giulia Fanti and Vyas Sekar

Grade: 4.0/4.0

Pittsburgh, PA, USA

2017–2023

Tsinghua University

Bachelor of Engineering, Department of Electronic Engineering

Grade: 92/100. Rank: 5/195

Beijing, China

2013–2017

Honors and Awards

CVPR Highlight, with Kaiyi Huang, Yukun Huang, Yu Li, Jianhong Bai, Xintao Wang, Xuefei Ning, Jiwen Yu, Pengfei Wan, Yu Wang, Xihui Liu 2026

2nd Place (Out of 2211 Teams) in AI Mathematical Olympiad 2, with Yichen You, Xuefei Ning, <https://www.kaggle.com/competitions/ai-mathematical-olympiad-progress-prize-2/leaderboard>. Solution: <https://github.com/imagination-research/aimo2/> 2025

Outstanding Reviewers (Top 5.6%) in CVPR 2025, <https://cvpr.thecvf.com/Conferences/2025/ProgramCommittee#all-outstanding-reviewer> 2025

12th Annual Best Scientific Cybersecurity Paper, with Boxin Wang, Weixin Chen, Hengzhi Pei, Chulin Xie, Mintong Kang, Chenhui Zhang, Chejian Xu, Zidi Xiong, Ritik Dutta, Rylan Schaeffer, Sang T. Truong, Simran Arora, Mantas Mazeika, Dan Hendrycks, Yu Cheng, Sanmi Koyejo, Dawn Song, Bo Li, granted by the National Security Agency (NSA), <https://www.nsa.gov/Press-Room/News-Highlights/Article/3990748/nsa-awards-authors-of-assessment-of-trustworthiness-in-gpt-models/> 2024

Best Machine Learning and Security Paper (Cybersecurity Best Paper Award) , with Boxin Wang, Weixin Chen, Hengzhi Pei, Chulin Xie, Mintong Kang, Chenhui Zhang, Chejian Xu, Zidi Xiong, Ritik Dutta, Rylan Schaeffer, Sang T. Truong, Simran Arora, Mantas Mazeika, Dan Hendrycks, Yu Cheng, Sanmi Koyejo, Dawn Song, Bo Li, granted by SpringerOpen, https://cybersecurity.springeropen.com/award-2024	2024
ICML Spotlight (Top 3.5%) , with Chulin Xie, Arturs Backurs, Sivakanth Gopi, Da Yu, Huseyin Inan, Harsha Nori, Haotian Jiang, Huishuai Zhang, Yin Tat Lee, Bo Li, Sergey Yekhanin	2024
NeurIPS Outstanding Paper & Oral (Top 0.045%) , with Boxin Wang, Weixin Chen, Hengzhi Pei, Chulin Xie, Mintong Kang, Chenhui Zhang, Chejian Xu, Zidi Xiong, Ritik Dutta, Rylan Schaeffer, Sang T. Truong, Simran Arora, Mantas Mazeika, Dan Hendrycks, Yu Cheng, Sanmi Koyejo, Dawn Song, Bo Li	2023
NeurIPS Oral (Synthetic Data Generation with Generative AI Workshop) , with Sivakanth Gopi, Janardhan Kulkarni, Harsha Nori, Sergey Yekhanin	2023
Top Reviewers (Top 2.6%) in NeurIPS 2023 , https://nips.cc/Conferences/2023/ProgramCommittee	2023
AAAI Scholarship , granted by Association for the Advancement of Artificial Intelligence	2022
Outstanding Reviewer (Top 8%) in NeurIPS 2021 , https://nips.cc/Conferences/2021/ProgramCommittee	2021
IMC Best Paper Finalist , with Alankar Jain, Chen Wang, Giulia Fanti, Vyas Sekar	2020
Top Reviewers (Top 7.6%) in ICML 2020 , https://icml.cc/Conferences/2020/Reviewers	2020
Cylab Presidential Fellowship , granted by Carnegie Mellon University	2020
Siemens FutureMakers Fellowship , granted by Siemens	2019
Best Reviewers (Top 400) in NeurIPS 2019 , https://nips.cc/Conferences/2019/Reviewers	2019
NeurIPS Spotlight (Top 4.1%) , with Kiran Thekumparampil, Ashish Khetan, and Sewoong Oh	2018
Presidential Fellowship , granted by Carnegie Mellon University	2017
Carnegie Institute of Technology Dean's Fellow , granted by Carnegie Mellon University	2017
Outstanding Bachelor Thesis , granted by Tsinghua University	2017
Meritorious Winner (9% Worldwide) , COMAP's Math Contest in Modeling	2015, 2016, 2017
National Scholarship , granted by the government of China	2014, 2015, 2016
Tsinghua Spark Class Fellowship (Top 1%) , for top students on scientific research	2015
The First Prize , National Physics Contest for College Student	2014

Experience

NVIDIA (Research Intern) <i>Host: Ming-Yu Liu, Xun Huang</i> ○ Topic: Denoising Diffusion Probabilistic Models (DDPM)	Santa Clara, CA, USA May 2021–Dec. 2021
Google (Research Intern) <i>Host: Yundi Qian</i> ○ Topic: compiler optimizations with reinforcement learning	Mountain View, CA, USA May 2020–Aug. 2020

Microsoft Research Asia (Research Intern)

Managers: Fei Gao, Taifeng Wang

- Topic: a large-scale empirical study of optimization methods

Beijing, China
Mar. 2017–Jun. 2017

Luogu Website (Cofounder and Developer)

<https://www.luogu.com.cn/>

- One of the biggest online judges in China.

China
2013

Skills

Programming Languages

C, C++, Python, Java, (Visual) Basic, Pascal, Haskell, MATLAB, Mathematica, PHP, JavaScript, HTML, CSS, SQL, Verilog, Assembly, bash, shell, L^AT_EX, etc.

Machine Learning Frameworks

PyTorch, TensorFlow, Theano, Keras, Blocks, CNTK, etc.

Teaching Assistant

CMU 18752: Estimation, Detection and Learning

Instructor: Rohit Negi

Pittsburgh, PA, USA
Spring 2020, Spring 2021

Services

Area Chair

Conference on Neural Information Processing Systems (NeurIPS) 2026

International Conference on Machine Learning (ICML) 2026

Reviewer

Conference on Neural Information Processing Systems (NeurIPS) 2019, 2020, 2021, 2022, 2023, 2025

International Conference on Machine Learning (ICML) 2020, 2021, 2022, 2023, 2024, 2025

International Conference on Learning Representations (ICLR) 2021, 2022, 2023, 2024, 2025, 2026

Computer Vision and Pattern Recognition Conference (CVPR) 2022, 2023, 2024, 2025, 2026

International Conference on Computer Vision (ICCV) 2023, 2025

European Conference on Computer Vision (ECCV) 2022, 2024

International Symposium on Information Theory (ISIT) 2020, 2021

Artificial Intelligence and Statistics (AISTATS) 2021

ICAIF Workshop on Synthetic Data for AI in Finance 2022

ACM CoNEXT NativeNI Workshop 2022

NeurIPS SyntheticData4ML Workshop 2022

Transactions on Machine Learning Research

Transactions on Pattern Analysis and Machine Intelligence

IET Image Processing
Transactions on Dependable and Secure Computing
Neural Networks
IEEE Access
IEEE/ACM Transactions on Networking
IEEE Transactions on Big Data
International Journal of Intelligent Systems

Publications (* denotes equal contribution, + denotes (co-)advising)

- [1] Thomas Humphries, **Zinan Lin**, and Sergey Yekhanin. “PE-means: Improved Differentially Private k -means Clustering through Private Evolution”. In: *arXiv preprint arXiv:2606.00342* (2026). URL: <https://arxiv.org/abs/2606.00342>.
- [2] Toan Tran, Arturs Backurs, **Zinan Lin**, Victor Reis, Li Xiong, and Sergey Yekhanin. “Differentially Private Synthetic Data via APIs 4: Tabular Data”. In: *International Conference on Machine Learning (ICML)* (2026). URL: <https://arxiv.org/abs/2606.08259>.
- [3] Chengquan Guo, Yuzhou Nie, Chulin Xie, **Zinan Lin**, Wenbo Guo, and Bo Li. “BlueCodeAgent: A Blue Teaming Agent Enabled by Automated Red Teaming for CodeGen AI”. In: *International Conference on Machine Learning (ICML)* (2026). URL: <https://arxiv.org/abs/2510.18131>.
- [4] Enshu Liu, Xuefei Ning, Yu Wang, and **Zinan Lin**+. “NI Sampling: Accelerating Discrete Diffusion Sampling by Token Order Optimization”. In: *International Conference on Learning Representations (ICLR)* (2026). URL: <https://openreview.net/forum?id=rrD1U0lzt5>.
- [5] Tianyuan Zou, **Zinan Lin**, Sivakanth Gopi, Yang Liu, Ya-Qin Zhang, Robert Sim, Xin Deng, and Sergey Yekhanin. “PE-SGD: Differentially Private Deep Learning via Evolution of Gradient Subspace for Text”. In: *International Conference on Learning Representations (ICLR)* (2026). URL: <https://openreview.net/forum?id=713ywmTZhV>.
- [6] Kaiyi Huang, Yukun Huang, Xintao Wang, **Zinan Lin**, Xuefei Ning, Pengfei Wan, Di Zhang, Yu Wang, and Xihui Liu. “FilMaster: Bridging Cinematic Principles and Generative AI for Automated Film Generation”. In: *International Conference on Learning Representations (ICLR)* (2026). URL: <https://arxiv.org/abs/2506.18899>.
- [7] Chengquan Guo, Chulin Xie, Yu Yang, Zhaorun Chen, **Zinan Lin**, Xander Davies, Yarin Gal, Dawn Song, and Bo Li. “RedCodeAgent: Automatic Red-teaming Agent against Diverse Code Agents”. In: *International Conference on Learning Representations (ICLR)* (2026). URL: <https://arxiv.org/abs/2510.02609>.
- [8] Enshu Liu, Theodore Zhengde Zhao, Xuefei Ning, **Zinan Lin**+, and Yu Wang. “NI Sampling++: Optimizing Token Order for Fast Discrete Diffusion Sampling with Reinforcement Learning”. In: *Conference on Language Modeling (COLM)* (2026). URL: <https://openreview.net/forum?id=i8zK3lfxYC>.

- [9] Chen Gong*, Kecen Li*, **Zinan Lin**, and Tianhao Wang. "From Easy to Hard++: Promoting Differentially Private Image Synthesis Through Spatial-Frequency Curriculum". In: *USENIX Security Symposium (USENIX Security)* (2026). URL: <https://arxiv.org/abs/2601.06368>.
- [10] Chen Gong, Kecen Li, **Zinan Lin**, and Tianhao Wang. "DP-SAPF: Saliency-Aware Parameter Fine-tuning of Public Models for Differentially Private Image Synthesis". In: *USENIX Security Symposium (USENIX Security)* (2026). URL: <https://arxiv.org/abs/2605.30312>.
- [11] Kaiyi Huang, Yukun Huang, Yu Li, Jianhong Bai, Xintao Wang, **Zinan Lin**, Xuefei Ning, Jiwen Yu, Pengfei Wan, Yu Wang, et al. "CineScene: Implicit 3D as Effective Scene Representation for Cinematic Video Generation". In: *The IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)* (2026). URL: <https://arxiv.org/abs/2602.06959>.
- [12] Kaiyi Huang, Yukun Huang, Xuefei Ning, **Zinan Lin**, Yu Wang, and Xihui Liu. "GenMAC: Compositional Text-to-Video Generation with Multi-Agent Collaboration". In: *AAAI Conference on Artificial Intelligence (AAAI)* (2026). URL: <https://arxiv.org/abs/2412.04440>.
- [13] Fangyuan Xu, Sihao Chen, **Zinan Lin**, Taiwei Shi, Sydney Graham, Pei Zhou, Mengting Wan, Alex Stein, Virginia Estellers, Charles Chen, Morris Sharp, Richard Speyer, Tadas Baltrusaitis, Jennifer Neville, Eunsol Choi, and Longqi Yang. "DP-RFT: Learning to Generate Synthetic Text via Differentially Private Reinforcement Fine-Tuning". In: *arXiv preprint arXiv:2602.18633* (2026). URL: <https://arxiv.org/abs/2602.18633>.
- [14] Yue Huang, Yu Jiang, Wenjie Wang, Haomin Zhuang, Xiaonan Luo, Yuchen Ma, Zhangchen Xu, Zichen Chen, Nuno Moniz, **Zinan Lin**, Pin-Yu Chen, Nitesh V Chawla, Nouha Dziri, Huan Sun, and Xiangliang Zhang. "Emergent Social Intelligence Risks in Generative Multi-Agent Systems". In: *ICML 2026 Workshop on Trustworthy AI for Good, ICML 2026 Workshop on Agents in the Wild: Safety, Security, and Beyond, ICML 2026 Workshop on Failure Modes in Agentic AI, ICML 2026 Workshop on Compositional Learning: Safety, Interpretability, and Agents* (2026). URL: <https://arxiv.org/abs/2603.27771>.
- [15] Shuaiqi Wang, Aadyaa Maddi, **Zinan Lin**, and Giulia Fanti. "ESDAE: Evaluating Synthetic Data for Agent Evaluation". In: *ICLR 2026 Workshop on AI with Recursive Self-Improvement, ICLR 2026 Workshop on Agents in the Wild: Safety, Security, and Beyond, ICLR 2026 Workshop on Navigating and Addressing Data Problems for Foundation Models* (2026). URL: <https://openreview.net/forum?id=yweA1eu5eP>.
- [16] Kecen Li, Chen Gong, Zinan Lin, Tianhao Wang, and Xiaokui Xiao. "Differentially Private Contrastive Learning via Bounding Group-level Contribution". In: *arXiv preprint arXiv:2604.26467* (2026). URL: <https://arxiv.org/abs/2604.26467>.
- [17] Zixin Ding, Junyuan Hong, Zhan Shi, Jiachen T Wang, **Zinan Lin**, Li Yin, Meng Liu, Zhangyang Wang, and Yuxin Chen. "Scaling Textual Gradients via Sampling-Based Momentum". In: *ACM Conference on AI and Agentic Systems (CAIS)* (2026). URL: <https://arxiv.org/abs/2506.00400>.
- [18] Shuaiqi Wang, Aadyaa Maddi, **Zinan Lin**, and Giulia Fanti. "SynAE: A Framework for Measuring the Quality of Synthetic Data for Tool-Calling Agent Evaluations". In: *arXiv preprint arXiv:2605.22564* (2026). URL: <https://arxiv.org/abs/2605.22564>.
- [19] Shuaiqi Wang, **Zinan Lin**, and Giulia Fanti. "Statistic Maximal Leakage". In: *Entropy* 28.7 (2026). ISSN: 1099-4300. DOI: [10.3390/e28070819](https://doi.org/10.3390/e28070819). URL: <https://www.mdpi.com/1099-4300/28/7/819>.

- [20] **Zinan Lin**, Enshu Liu, Xuefei Ning, Junyi Zhu, Wenyu Wang, and Sergey Yekhanin. “Latent Zoning Network: A Unified Principle for Generative Modeling, Representation Learning, and Classification”. In: *Advances in Neural Information Processing Systems (NeurIPS)* (2025). URL: <https://arxiv.org/abs/2509.15591>.
- [21] Shuaiqi Wang*, Vikas Raunak*, Arturs Backurs, Victor Reis, Pei Zhou+, Sihao Chen, Longqi Yang, **Zinan Lin+**, Sergey Yekhanin, and Giulia Fanti. “Struct-Bench: A Benchmark for Differentially Private Structured Text Generation”. In: *Advances in Neural Information Processing Systems (NeurIPS)* (2025). URL: <https://arxiv.org/abs/2509.10696>.
- [22] Haoxiang Wang, **Zinan Lin**, Da Yu, and Huishuai Zhang. “Synthesize Privacy-Preserving High-Resolution Images via Private Textual Intermediaries”. In: *Advances in Neural Information Processing Systems (NeurIPS)* (2025). URL: <https://arxiv.org/abs/2506.07555>.
- [23] Enshu Liu, Qian Chen, Xuefei Ning, Shengen Yan, Guohao Dai, **Zinan Lin+**, and Yu Wang. “Distilled Decoding 2: One-step Sampling of Image Auto-regressive Models with Conditional Score Distillation”. In: *Advances in Neural Information Processing Systems (NeurIPS)* (2025). URL: <https://arxiv.org/abs/2510.21003>.
- [24] **Zinan Lin**, Tadas Baltrusaitis, Wenyu Wang, and Sergey Yekhanin. “Differentially Private Synthetic Data via APIs 3: Using Simulators Instead of Foundation Models”. In: *ICLR Workshop on Data Problems, ICLR Workshop on Synthetic Data* (2025). URL: <https://arxiv.org/abs/2502.05505>.
- [25] Enshu Liu, Xuefei Ning, Yu Wang, and **Zinan Lin+**. “Distilled Decoding 1: One-step Sampling of Image Auto-regressive Models with Flow Matching”. In: *International Conference on Learning Representations (ICLR)* (2025). URL: <https://arxiv.org/abs/2412.17153>.
- [26] Enshu Liu*, Junyi Zhu*, **Zinan Lin+**, Xuefei Ning+, Matthew B Blaschko, Sergey Yekhanin, Shengen Yan, Guohao Dai, Huazhong Yang, and Yu Wang. “Linear Combination of Saved Checkpoints Makes Consistency and Diffusion Models Better”. In: *International Conference on Learning Representations (ICLR)* (2025). URL: <https://arxiv.org/abs/2404.02241>.
- [27] Chejian Xu, Jiawei Zhang, Zhaorun Chen, Chulin Xie, Mintong Kang, Zhuowen Yuan, Zidi Xiong, Chenhui Zhang, Lingzhi Yuan, Yi Zeng, Peiyang Xu, Chengquan Guo, Andy Zhou, Jeffrey Ziwei Tan, Zhun Wang, Alexander Xiong, Xuandong Zhao, Yu Gai, Francesco Pinto, Yujin Potter, Zhen Xiang, **Zinan Lin**, Dan Hendrycks, Dawn Song, and Bo Li. “MMDT: Decoding the Trustworthiness and Safety of Multimodal Foundation Models”. In: *International Conference on Learning Representations (ICLR)* (2025). URL: <https://arxiv.org/abs/2503.14827>.
- [28] Tianchen Zhao, Tongcheng Fang, Enshu Liu, Wan Rui, Widyadewi Soedarmadji, Shiyao Li, **Zinan Lin**, Guohao Dai, Shengen Yan, Huazhong Yang, Xuefei Ning, and Yu Wang. “ViDiT-Q: Efficient and Accurate Quantization of Diffusion Transformers for Image and Video Generation”. In: *International Conference on Learning Representations (ICLR)* (2025). URL: <https://arxiv.org/abs/2406.02540>.
- [29] Gong Chen*, Kecen Li*, **Zinan Lin**, and Wang Tianhao. “DPIImageBench: A Unified Benchmark for Differentially Private Image Synthesis”. In: *The ACM Conference on Computer and Communications Security (CCS)* (2025). URL: <https://arxiv.org/abs/2503.14681>.

- [30] Shuaiqi Wang, Shuran Zheng, **Zinan Lin**, Giulia Fanti, and Zhiwei Steven Wu. “Inferentially-Private Private Information”. In: *The ACM Web Conference (WWW)* (2025). URL: <https://arxiv.org/abs/2410.17095>.
- [31] YiPeng Chen, Chulin Xie, **Zinan Lin**, Dawn Song, and Bo Li. “KnowData: Knowledge-Enabled Data Generation for Improving Multimodal Models”. In: (2025). URL: <https://openreview.net/forum?id=FqWtMGw8tt>.
- [32] Xuefei Ning*, Zifu Wang*, Shiyao Li*, **Zinan Lin***, Peiran Yao*, Tianyu Fu, Matthew B Blaschko, Guohao Dai, Huazhong Yang, and Yu Wang. “Can LLMs Learn by Teaching for Better Reasoning? A Preliminary Study”. In: *Advances in Neural Information Processing Systems (NeurIPS)* (2024). URL: <https://arxiv.org/abs/2406.14629>.
- [33] Sangyun Lee, **Zinan Lin**, and Giulia Fanti. “Improving the Training of Rectified Flows”. In: *Advances in Neural Information Processing Systems (NeurIPS)* (2024). URL: <https://arxiv.org/abs/2405.20320>.
- [34] Chengquan Guo*, Xun Liu*, Chulin Xie*, Andy Zhou, Yi Zeng, **Zinan Lin**, Dawn Song, and Bo Li. “RedCode: Risky Code Execution and Generation Benchmark for Code Agents”. In: *Advances in Neural Information Processing Systems (NeurIPS)* (2024). URL: <https://arxiv.org/abs/2411.07781>.
- [35] Enshu Liu*, Junyi Zhu*, **Zinan Lin+**, Xuefei Ning+, Matthew B Blaschko, Shengen Yan, Guohao Dai, Huazhong Yang, and Yu Wang. “Efficient Expert Pruning for Sparse Mixture-of-Experts Language Models: Enhancing Performance and Reducing Inference Costs”. In: *arXiv preprint arXiv:2407.00945* (2024). URL: <https://arxiv.org/abs/2407.00945>.
- [36] **Zinan Lin**, Sivakanth Gopi, Janardhan Kulkarni, Harsha Nori, and Sergey Yekhanin. “Differentially Private Synthetic Data via Foundation Model APIs 1: Images”. In: *International Conference on Learning Representations (ICLR)* (2024). URL: <https://arxiv.org/abs/2305.15560>.
- [37] Chulin Xie, **Zinan Lin**, Arturs Backurs, Sivakanth Gopi, Da Yu, Huseyin Inan, Harsha Nori, Haotian Jiang, Huishuai Zhang, Yin Tat Lee, Bo Li, and Sergey Yekhanin. “Differentially Private Synthetic Data via Foundation Model APIs 2: Text”. In: *International Conference on Machine Learning (ICML)* (2024). URL: <https://arxiv.org/abs/2403.01749>.
- [38] Xuefei Ning*, **Zinan Lin***, Zixuan Zhou*, Huazhong Yang, and Yu Wang. “Skeleton-of-Thought: Prompting LLMs for Efficient Parallel Generation”. In: *International Conference on Learning Representations (ICLR)* (2024). URL: <https://arxiv.org/abs/2307.15337>.
- [39] Shuaiqi Wang, **Zinan Lin**, and Giulia Fanti. “Statistic Maximal Leakage”. In: *The IEEE International Symposium on Information Theory (ISIT)* (2024). URL: <https://arxiv.org/abs/2411.18531>.
- [40] Xinyu Tang, Richard Shin, Huseyin A. Inan, Andre Manoel, Fatemehsadat Miresghallah, **Zinan Lin**, Sivakanth Gopi, Janardhan Kulkarni, and Robert Sim. “Privacy-Preserving In-Context Learning with Differentially Private Few-Shot Generation”. In: *International Conference on Learning Representations (ICLR)* (2024). URL: <https://arxiv.org/abs/2309.11765>.
- [41] Arturs Backurs, **Zinan Lin**, Sepideh Mahabadi, Sandeep Silwal, and Jakub Tarnawski. “Efficiently Computing Similarities to Private Datasets”. In: *International Conference on Learning Representations (ICLR)* (2024). URL: <https://arxiv.org/abs/2403.08917>.

- [42] Junyi Zhu, **Zinan Lin**, Enshu Liu, Xuefei Ning, and Matthew B. Blaschko. “Rescaling Intermediate Features Makes Trained Consistency Models Perform Better”. In: *International Conference on Learning Representations (ICLR) (TinyPapers)* (2024). URL: <https://openreview.net/forum?id=1o3LnwflAl>.
- [43] Ronghao Ni, **Zinan Lin**, Shuaiqi Wang, and Giulia Fanti. “Mixture-of-Linear-Experts for Long-term Time Series Forecasting”. In: *International Conference on Artificial Intelligence and Statistics (AISTATS)* (2024). URL: <https://arxiv.org/abs/2312.06786>.
- [44] Lin Zhao*, Tianchen Zhao*, **Zinan Lin**, Xuefei Ning, Guohao Dai, Huazhong Yang, and Yu Wang. “FlashEval: Towards Fast and Accurate Evaluation of Text-to-image Diffusion Generative Models”. In: *The IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)* (2024). URL: <https://arxiv.org/abs/2403.16379>.
- [45] **Zinan Lin***, Shuaiqi Wang*, Vyas Sekar, and Giulia Fanti. “Summary Statistic Privacy in Data Sharing”. In: *IEEE Journal on Selected Areas in Information Theory (JSAIT)* (2024). URL: <https://arxiv.org/abs/2303.02014>.
- [46] Da Yu, Sivakanth Gopi, Janardhan Kulkarni, **Zinan Lin**, Saurabh Naik, Tomasz Lukasz Religa, Jian Yin, and Huishuai Zhang. “Selective Pre-training for Private Fine-tuning”. In: *Transactions on Machine Learning Research (TMLR)* (2024). URL: <https://arxiv.org/abs/2305.13865>.
- [47] Tianchen Zhao, Xuefei Ning, Tongcheng Fang, Enshu Liu, Guyue Huang, **Zinan Lin**, Shengen Yan, Guohao Dai, and Yu Wang. “MixDQ: Memory-Efficient Few-Step Text-to-Image Diffusion Models with Metric-Decoupled Mixed Precision Quantization”. In: *European Conference on Computer Vision (ECCV)* (2024). URL: <https://arxiv.org/abs/2405.17873>.
- [48] Enshu Liu*, Xuefei Ning*, **Zinan Lin***, Huazhong Yang, and Yu Wang. “OMS-DPM: Optimizing the Model Schedule for Diffusion Probabilistic Models”. In: *International Conference on Machine Learning (ICML)*. 2023. URL: <https://arxiv.org/abs/2306.08860>.
- [49] Boxin Wang, Weixin Chen, Hengzhi Pei, Chulin Xie, Mintong Kang, Chenhui Zhang, Chejian Xu, Zidi Xiong, Ritik Dutta, Rylan Schaeffer, Sang T Truong, Simran Arora, Mantas Mazeika, Dan Hendrycks, **Zinan Lin**, Yu Cheng, Sanmi Koyejo, Dawn Song, and Bo Li. “DecodingTrust: A Comprehensive Assessment of Trustworthiness in GPT Models”. In: *Advances in Neural Information Processing Systems (NeurIPS)*. 2023. URL: <https://arxiv.org/abs/2306.11698>.
- [50] Da Yu, Arturs Backurs, Sivakanth Gopi, Huseyin Inan, Janardhan Kulkarni, **Zinan Lin**, Chulin Xie, Huishuai Zhang, and Wanrong Zhang. “Training Private and Efficient Language Models with Synthetic Data from LLMs”. In: *NeurIPS Workshop on Socially Responsible Language Modelling Research* (2023). URL: <https://openreview.net/forum?id=FKwtKzglFb>.
- [51] **Zinan Lin**. “Data Sharing with Generative Adversarial Networks: From Theory to Practice”. PhD thesis. Carnegie Mellon University, 2022. URL: https://zinanlin.me/bio/PhD_thesis.pdf.
- [52] **Zinan Lin***, Shuaiqi Wang*, Vyas Sekar, and Giulia Fanti. “Distributional Privacy for Data Sharing”. In: *NeurIPS 2022 Workshop on Synthetic Data for Empowering ML Research*. 2022. URL: <https://openreview.net/forum?id=6oVAzFsHlFK>.
- [53] Yucheng Yin, **Zinan Lin**, Minhao Jin, Giulia Fanti, and Vyas Sekar. “Practical GAN-based Synthetic IP Header Trace Generation using NetShare”. In: *ACM Special Interest Group on Data Communication (SIGCOMM)*. 2022. URL: <https://dl.acm.org/doi/10.1145/3544216.3544251>.

- [54] **Zinan Lin**, Hao Liang, Giulia Fanti, and Vyas Sekar. “RareGAN: Generating Samples for Rare Classes”. In: *AAAI Conference on Artificial Intelligence (AAAI)*. 2022. URL: <https://ojs.aaai.org/index.php/AAAI/article/view/20715>.
- [55] Yucheng Yin, **Zinan Lin**, Minhao Jin, Giulia Fanti, and Vyas Sekar. “PcapShare: Exploring the Feasibility of GANs for Synthetic Packet Header Trace Generation”. In: *International Conference on COMMunication Systems and NETworkS (COMSNETS) (demo)*. 2022. URL: https://www.comsnets.org/demos_exhibits.html.
- [56] **Zinan Lin**, Vyas Sekar, and Giulia Fanti. “Why Spectral Normalization Stabilizes GANs: Analysis and Improvements”. In: *Advances in Neural Information Processing Systems (NeurIPS)*. 2021. URL: <http://arxiv.org/abs/2009.02773>.
- [57] **Zinan Lin**, Vyas Sekar, and Giulia Fanti. “On the Privacy Properties of GAN-generated Samples”. In: *International Conference on Artificial Intelligence and Statistics (AISTATS)*. PMLR. 2021, pp. 1522–1530. URL: <https://arxiv.org/abs/2206.01349>.
- [58] Todd Huster, Jeremy E.J. Cohen, **Zinan Lin**, Kevin Chan, Cho-Yu Jason Chiang, and Vyas Sekar. “Pareto GAN: Extending the Representational Power of GANs to Heavy-Tailed Distributions”. In: *International Conference on Machine Learning (ICML)*. 2021. URL: <http://proceedings.mlr.press/v139/huster21a.html>.
- [59] Mircea Trofin, Yundi Qian, Eugene Brevdo, **Zinan Lin**, Krzysztof Choromanski, and David Li. “MLGO: a Machine Learning Guided Compiler Optimizations Framework”. In: *arXiv preprint arXiv:2101.04808* (2021). URL: <https://arxiv.org/abs/2101.04808>.
- [60] **Zinan Lin**, Kiran Koshy Thekumparampil, Giulia Fanti, and Sewoong Oh. “InfoGAN-CR and ModelCentrality: Self-supervised Model Training and Selection for Disentangling GANs”. In: *International Conference on Machine Learning (ICML)*. 2020, pp. 7775–7786. URL: <https://arxiv.org/abs/1906.06034>.
- [61] **Zinan Lin**, Alankar Jain, Chen Wang, Giulia Fanti, and Vyas Sekar. “Using GANs for Sharing Networked Timeseries Data: Challenges, Initial Promise, and Open Questions”. In: *Proceedings of the Internet Measurement Conference (IMC)*. 2020. URL: <http://arxiv.org/abs/1909.13403>.
- [62] **Zinan Lin**, Ashish Khetan, Giulia Fanti, and Sewoong Oh. “PacGAN: The Power of Two Samples in Generative Adversarial Networks”. In: *IEEE Journal on Selected Areas in Information Theory (JSAIT)* 1.1 (2020), pp. 324–335. URL: <https://ieeexplore.ieee.org/document/9046238>.
- [63] **Zinan Lin**, Soo-Jin Moon, Carolina M. Zarate, Ritika Mulagalapalli, Sekar Kulandaivel, Giulia Fanti, and Vyas Sekar. “Towards Oblivious Network Analysis using Generative Adversarial Networks”. In: *Proceedings of the 18th ACM Workshop on Hot Topics in Networks (HotNets)*. ACM. 2019. URL: <https://dl.acm.org/doi/10.1145/3365609.3365854>.
- [64] **Zinan Lin**, Ashish Khetan, Giulia Fanti, and Sewoong Oh. “PacGAN: The Power of Two Samples in Generative Adversarial Networks”. In: *Advances in Neural Information Processing Systems (NeurIPS)*. 2018, pp. 1498–1507. URL: <https://arxiv.org/abs/1712.04086>.
- [65] Kiran K Thekumparampil, Ashish Khetan, **Zinan Lin**, and Sewoong Oh. “Robustness of Conditional GANs to Noisy Labels”. In: *Advances in Neural Information Processing Systems (NeurIPS)*. 2018, pp. 10271–10282. URL: <https://arxiv.org/abs/1811.03205>.

- [66] **Zinan Lin**, Yongfeng Huang, and Jilong Wang. “RNN-SM: Fast Steganalysis of VoIP Streams Using Recurrent Neural Network”. In: *IEEE Transactions on Information Forensics and Security (TIFS)* 13.7 (July 2018), pp. 1854–1868. ISSN: 1556-6013. DOI: [10.1109/TIFS.2018.2806741](https://doi.org/10.1109/TIFS.2018.2806741). URL: <http://ieeexplore.ieee.org/document/8292900>.