

Zinan Lin

Microsoft Building 99, 14820 NE 36th St – Redmond, WA 98052 – USA

✉ zinanlin@microsoft.com • 🌐 zinanlin.me

🔍 scholar.google.com/citations?user=67nE-wQ_g_cC

🔗 github.com/fjxmlzn

Employment

Microsoft Research

Senior Researcher

Redmond, WA, USA

Oct. 2022-Present

Education

Carnegie Mellon University

Ph.D., Department of Electrical and Computer Engineering

Secondary Master in Machine Learning, School of Computer Science

Secondary Master in Electrical and Computer Engineering

Ph.D. advisors: Giulia Fanti and Vyas Sekar

Grade: 4.0/4.0

Pittsburgh, PA, USA

2017–2023

Tsinghua University

Bachelor of Engineering, Department of Electronic Engineering

Grade: 92/100. Rank: 5/195

Beijing, China

2013–2017

Honors and Awards

NeurIPS Outstanding Paper & Oral, with Boxin Wang, Weixin Chen, Hengzhi Pei, Chulin Xie, Mintong Kang, Chenhui Zhang, Chejian Xu, Zidi Xiong, Ritik Dutta, Rylan Schaeffer, Sang T. Truong, Simran Arora, Mantas Mazeika, Dan Hendrycks, Yu Cheng, Sanmi Koyejo, Dawn Song, Bo Li 2023

NeurIPS Oral (Synthetic Data Generation with Generative AI Workshop), with Sivakanth Gopi, Janardhan Kulkarni, Harsha Nori, Sergey Yekhanin 2023

Top Reviewers (Top 2.6%) in NeurIPS 2023, 2023
<https://nips.cc/Conferences/2023/ProgramCommittee>

AAAI Scholarship, granted by Association for the Advancement of Artificial Intelligence 2022

Outstanding Reviewer (Top 8%) in NeurIPS 2021, 2021
<https://nips.cc/Conferences/2021/ProgramCommittee>

IMC Best Paper Finalist, with Alankar Jain, Chen Wang, Giulia Fanti, Vyas Sekar 2020

Top Reviewers (Top 7.6%) in ICML 2020, <https://icml.cc/Conferences/2020/Reviewers> 2020

Cylab Presidential Fellowship, granted by Carnegie Mellon University 2020

Siemens FutureMakers Fellowship, granted by Siemens 2019

Best Reviewers (Top 400) in NeurIPS 2019, <https://nips.cc/Conferences/2019/Reviewers> 2019

NeurIPS Spotlight , with Kiran Thekumparampil, Ashish Khetan, and Sewoong Oh	2018
Presidential Fellowship , granted by Carnegie Mellon University	2017
Carnegie Institute of Technology Dean's Fellow , granted by Carnegie Mellon University	2017
Outstanding Bachelor Thesis , granted by Tsinghua University	2017
Meritorious Winner (9% Worldwide) , COMAP's Math Contest in Modeling	2015, 2016, 2017
National Scholarship , granted by the government of China	2014, 2015, 2016
Tsinghua Spark Class Fellowship (Top 1%) , for top students on scientific research	2015
The First Prize , National Physics Contest for College Student	2014

Experience

NVIDIA (Research Intern) <i>Host: Ming-Yu Liu, Xun Huang</i> ○ Topic: Denoising Diffusion Probabilistic Models (DDPM)	Santa Clara, CA, USA May 2021–Dec. 2021
Google (Research Intern) <i>Host: Yundi Qian</i> ○ Topic: compiler optimizations with reinforcement learning	Mountain View, CA, USA May 2020–Aug. 2020
Microsoft Research Asia (Research Intern) <i>Managers: Fei Gao, Taifeng Wang</i> ○ Topic: a large-scale empirical study of optimization methods	Beijing, China Mar. 2017–Jun. 2017
Luogu Website (Cofounder and Developer) <i>https://www.luogu.com.cn/</i> ○ One of the biggest online judges in China.	China 2013

Skills

Programming Languages

C, C++, Python, Java, (Visual) Basic, Pascal, Haskell, MATLAB, Mathematica, PHP, JavaScript, HTML, CSS, SQL, Verilog, Assembly, bash, shell, L^AT_EX, etc.

Machine Learning Frameworks

TensorFlow, PyTorch, Theano, Keras, Blocks, CNTK, etc.

Teaching Assistant

CMU 18752: Estimation, Detection and Learning <i>Instructor: Rohit Negi</i>	Pittsburgh, PA, USA Spring 2020, Spring 2021
---	--

Services

I am serving as a reviewer/PC member for:

Conference on Neural Information Processing Systems (NeurIPS)	2019, 2020, 2021, 2022, 2023
International Conference on Machine Learning (ICML)	2020, 2021, 2022, 2023, 2024

International Conference on Learning Representations (ICLR)	2021, 2022, 2023, 2024
Computer Vision and Pattern Recognition Conference (CVPR)	2022, 2023, 2024
International Conference on Computer Vision (ICCV)	2023
European Conference on Computer Vision (ECCV)	2022, 2024
International Symposium on Information Theory (ISIT)	2020, 2021
Artificial Intelligence and Statistics (AISTATS)	2021
ICAIF Workshop on Synthetic Data for AI in Finance	2022
ACM CoNEXT NativeNI Workshop	2022
NeurIPS SyntheticData4ML Workshop	2022
Transactions on Machine Learning Research	
Transactions on Pattern Analysis and Machine Intelligence	
IET Image Processing	
Transactions on Dependable and Secure Computing	
Neural Networks	
IEEE Access	
IEEE/ACM Transactions on Networking	
IEEE Transactions on Big Data	
International Journal of Intelligent Systems	

Publications (* denotes equal contribution, + denotes co-advising)

- [1] Enshu Liu*, Junyi Zhu*, **Zinan Lin**+, Xuefei Ning+, Matthew B Blaschko, Sergey Yekhanin, Shengen Yan, Guohao Dai, Huazhong Yang, and Yu Wang. "Linear Combination of Saved Checkpoints Makes Consistency and Diffusion Models Better". In: *arXiv preprint arXiv:2404.02241* (2024). URL: <https://arxiv.org/abs/2404.02241>.
- [2] **Zinan Lin**, Sivakanth Gopi, Janardhan Kulkarni, Harsha Nori, and Sergey Yekhanin. "Differentially Private Synthetic Data via Foundation Model APIs 1: Images". In: *International Conference on Learning Representations (ICLR)* (2024). URL: <https://arxiv.org/abs/2305.15560>.
- [3] Chulin Xie, **Zinan Lin**, Arturs Backurs, Sivakanth Gopi, Da Yu, Huseyin Inan, Harsha Nori, Haotian Jiang, Huishuai Zhang, Yin Tat Lee, Bo Li, and Sergey Yekhanin. "Differentially Private Synthetic Data via Foundation Model APIs 2: Text". In: *arXiv preprint arXiv:2403.01749* (2024). URL: <https://arxiv.org/abs/2403.01749>.
- [4] Xuefei Ning*, **Zinan Lin***, Zixuan Zhou*, Huazhong Yang, and Yu Wang. "Skeleton-of-Thought: Large Language Models Can Do Parallel Decoding". In: *International Conference on Learning Representations (ICLR)* (2024). URL: <https://arxiv.org/abs/2307.15337>.
- [5] Shuaiqi Wang, **Zinan Lin**, and Giulia Fanti. "Statistic Maximal Leakage". In: *The IEEE International Symposium on Information Theory (ISIT)* (2024).
- [6] Xinyu Tang, Richard Shin, Huseyin A. Inan, Andre Manoel, Fatemehsadat Miresghallah, **Zinan Lin**, Sivakanth Gopi, Janardhan Kulkarni, and Robert Sim. "Privacy-Preserving In-Context

- Learning with Differentially Private Few-Shot Generation". In: *International Conference on Learning Representations (ICLR)* (2024). URL: <https://arxiv.org/abs/2309.11765>.
- [7] Arturs Backurs, **Zinan Lin**, Sepideh Mahabadi, Sandeep Silwal, and Jakub Tarnawski. "Efficiently Computing Similarities to Private Datasets". In: *International Conference on Learning Representations (ICLR)* (2024). URL: <https://arxiv.org/abs/2403.08917>.
 - [8] Junyi Zhu, **Zinan Lin**, Enshu Liu, Xuefei Ning, and Matthew B. Blaschko. "Rescaling Intermediate Features Makes Trained Consistency Models Perform Better". In: *International Conference on Learning Representations (ICLR) (TinyPapers)* (2024). URL: <https://openreview.net/forum?id=1o3LnwflAl>.
 - [9] Ronghao Ni, **Zinan Lin**, Shuaiqi Wang, and Giulia Fanti. "Mixture-of-Linear-Experts for Long-term Time Series Forecasting". In: *International Conference on Artificial Intelligence and Statistics (AISTATS)* (2024). URL: <https://arxiv.org/abs/2312.06786>.
 - [10] Lin Zhao*, Tianchen Zhao*, **Zinan Lin**, Xuefei Ning, Guohao Dai, Huazhong Yang, and Yu Wang. "FlashEval: Towards Fast and Accurate Evaluation of Text-to-image Diffusion Generative Models". In: *The IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)* (2024). URL: <https://arxiv.org/abs/2403.16379>.
 - [11] Enshu Liu*, Xuefei Ning*, **Zinan Lin***, Huazhong Yang, and Yu Wang. "OMS-DPM: Optimizing the Model Schedule for Diffusion Probabilistic Models". In: *Proceedings of Machine Learning and Systems (ICML)*. 2023. URL: <https://arxiv.org/abs/2306.08860>.
 - [12] Boxin Wang, Weixin Chen, Hengzhi Pei, Chulin Xie, Mintong Kang, Chenhui Zhang, Chejian Xu, Zidi Xiong, Ritik Dutta, Rylan Schaeffer, Sang T Truong, Simran Arora, Mantas Mazeika, Dan Hendrycks, **Zinan Lin**, Yu Cheng, Sanmi Koyejo, Dawn Song, and Bo Li. "DecodingTrust: A Comprehensive Assessment of Trustworthiness in GPT Models". In: *Advances in Neural Information Processing Systems (NeurIPS)*. 2023. URL: <https://arxiv.org/abs/2306.11698>.
 - [13] **Zinan Lin***, Shuaiqi Wang*, Vyas Sekar, and Giulia Fanti. "Summary Statistic Privacy in Data Sharing". In: *arXiv preprint arXiv:2303.02014* (2023). URL: <https://arxiv.org/abs/2303.02014>.
 - [14] Da Yu, Arturs Backurs, Sivakanth Gopi, Huseyin Inan, Janardhan Kulkarni, **Zinan Lin**, Chulin Xie, Huishuai Zhang, and Wanrong Zhang. "Training Private and Efficient Language Models with Synthetic Data from LLMs". In: *NeurIPS Workshop on Socially Responsible Language Modelling Research* (2023).
 - [15] Da Yu, Sivakanth Gopi, Janardhan Kulkarni, **Zinan Lin**, Saurabh Naik, Tomasz Lukasz Religa, Jian Yin, and Huishuai Zhang. "Selective Pre-training for Private Fine-tuning". In: *arXiv preprint arXiv:2305.13865* (2023). URL: <https://arxiv.org/abs/2305.13865>.
 - [16] **Zinan Lin**. "Data Sharing with Generative Adversarial Networks: From Theory to Practice". PhD thesis. Carnegie Mellon University, 2022. URL: https://zinanlin.me/bio/PhD_thesis.pdf.
 - [17] **Zinan Lin***, Shuaiqi Wang*, Vyas Sekar, and Giulia Fanti. "Distributional Privacy for Data Sharing". In: *NeurIPS 2022 Workshop on Synthetic Data for Empowering ML Research*. URL: <https://openreview.net/forum?id=6oVAzFsHlFK>.

- [18] Yucheng Yin, **Zinan Lin**, Minhao Jin, Giulia Fanti, and Vyas Sekar. “Practical GAN-based Synthetic IP Header Trace Generation using NetShare”. In: *ACM Special Interest Group on Data Communication (SIGCOMM)*. 2022. URL: <https://dl.acm.org/doi/10.1145/3544216.3544251>.
- [19] **Zinan Lin**, Hao Liang, Giulia Fanti, and Vyas Sekar. “RareGAN: Generating Samples for Rare Classes”. In: *AAAI Conference on Artificial Intelligence (AAAI)*. 2022. URL: <https://ojs.aaai.org/index.php/AAAI/article/view/20715>.
- [20] Yucheng Yin, **Zinan Lin**, Minhao Jin, Giulia Fanti, and Vyas Sekar. “PcapShare: Exploring the Feasibility of GANs for Synthetic Packet Header Trace Generation”. In: *International Conference on COMMunication Systems and NETWORKS (COMSNETS) (demo)*. 2022. URL: https://www.comsnets.org/demos_exhibits.html.
- [21] **Zinan Lin**, Vyas Sekar, and Giulia Fanti. “Why Spectral Normalization Stabilizes GANs: Analysis and Improvements”. In: *Advances in Neural Information Processing Systems (NeurIPS)*. 2021. URL: <http://arxiv.org/abs/2009.02773>.
- [22] **Zinan Lin**, Vyas Sekar, and Giulia Fanti. “On the Privacy Properties of GAN-generated Samples”. In: *International Conference on Artificial Intelligence and Statistics (AISTATS)*. PMLR. 2021, pp. 1522–1530. URL: <https://arxiv.org/abs/2206.01349>.
- [23] Todd Huster, Jeremy E.J. Cohen, **Zinan Lin**, Kevin Chan, Cho-Yu Jason Chiang, and Vyas Sekar. “Pareto GAN: Extending the Representational Power of GANs to Heavy-Tailed Distributions”. In: *Proceedings of Machine Learning and Systems (ICML)*. 2021. URL: <http://proceedings.mlr.press/v139/huster21a.html>.
- [24] Mircea Trofin, Yundi Qian, Eugene Brevdo, **Zinan Lin**, Krzysztof Choromanski, and David Li. “MLGO: a Machine Learning Guided Compiler Optimizations Framework”. In: *arXiv preprint arXiv:2101.04808* (2021). URL: <https://arxiv.org/abs/2101.04808>.
- [25] **Zinan Lin**, Kiran Koshy Thekumparampil, Giulia Fanti, and Sewoong Oh. “InfoGAN-CR and ModelCentrality: Self-supervised Model Training and Selection for Disentangling GANs”. In: *Proceedings of Machine Learning and Systems (ICML)*. 2020, pp. 7775–7786. URL: <https://arxiv.org/abs/1906.06034>.
- [26] **Zinan Lin**, Alankar Jain, Chen Wang, Giulia Fanti, and Vyas Sekar. “Using GANs for Sharing Networked Timeseries Data: Challenges, Initial Promise, and Open Questions”. In: *Proceedings of the Internet Measurement Conference (IMC)*. 2020. URL: <http://arxiv.org/abs/1909.13403>.
- [27] **Zinan Lin**, Ashish Khetan, Giulia Fanti, and Sewoong Oh. “PacGAN: The Power of Two Samples in Generative Adversarial Networks”. In: *IEEE Journal on Selected Areas in Information Theory (JSAIT)* 1.1 (2020), pp. 324–335. URL: <https://ieeexplore.ieee.org/document/9046238>.
- [28] **Zinan Lin**, Soo-Jin Moon, Carolina M. Zarate, Ritika Mulagalapalli, Sekar Kulandaivel, Giulia Fanti, and Vyas Sekar. “Towards Oblivious Network Analysis using Generative Adversarial Networks”. In: *Proceedings of the 18th ACM Workshop on Hot Topics in Networks (HotNets)*. ACM. 2019. URL: <https://dl.acm.org/doi/10.1145/3365609.3365854>.
- [29] **Zinan Lin**, Ashish Khetan, Giulia Fanti, and Sewoong Oh. “PacGAN: The Power of Two Samples in Generative Adversarial Networks”. In: *Advances in Neural Information Processing Systems (NeurIPS)*. 2018, pp. 1498–1507. URL: <https://arxiv.org/abs/1712.04086>.

- [30] Kiran K Thekumparampil, Ashish Khetan, **Zinan Lin**, and Sewoong Oh. “Robustness of Conditional GANs to Noisy Labels”. In: *Advances in Neural Information Processing Systems (NeurIPS)*. 2018, pp. 10271–10282. URL: <https://arxiv.org/abs/1811.03205>.
- [31] **Zinan Lin**, Yongfeng Huang, and Jilong Wang. “RNN-SM: Fast Steganalysis of VoIP Streams Using Recurrent Neural Network”. In: *IEEE Transactions on Information Forensics and Security (TIFS)* 13.7 (July 2018), pp. 1854–1868. ISSN: 1556-6013. DOI: 10.1109/TIFS.2018.2806741. URL: <http://ieeexplore.ieee.org/document/8292900>.